

Part 1. Search

relax agent: based on current state; planning agent: consequence of actions choose best!
A search problem is consisted of ① A state space, ② A set of actions

③ a transition model ④ An action cost ⑤ Start state ⑥ Goal State

property	Optimal Tree Graph	Complete Tree Graph	Time	Space
BFS	FIFO queue Cost=1	✓	✓	b ^s
DFS	LIFO Stack each time search depth = 1, 2, ... uniform cost search	X	X	✓
Iterative Deepening	Cost=1	X	b ^s	b ^m
A* Search	fun=g+heur fun=g+h+heur	✓	X	

⑦ admissibility: $\forall n \ 0 \leq h(n) \leq h^*(n)$

⑧ consistency: $\forall A, C, H(A) - h_A \leq Cost(A, C)$.

⑨ dominance: $\forall n \ h(n) \geq h^*(n)$.

Theorem 1: if h is admissible, then A^* yields an optimal solution in tree search.

Theorem 2: if h is consistent, then A^* yields an optimal solution in graph search.

greedy search: select the one with lowest heuristic value to expand.

not complete & optimal.

Local search:

→ not L.O.

① Hill climbing search: best neighbor

② Simulated Annealing:

```
function SIMULATED-ANNEALING(problem,schedule) returns a state
    current ← problem.initial-state
    for t = 1 to ∞ do
```

T ← schedule(t)

if T = 0 then return current

next ← a randomly selected successor of current

ΔE ← next.value - current.value

if ΔE > 0 then current ← next P = 1

ΔE ≤ 0 else current ← next only with probability e^{ΔE/T}



③ Local Beam search: best K result each time

④ Genetic Algorithms:

```
function GENETIC-ALGORITHM(population, FITNESS-FN) returns an individual
    inputs: population, a set of individuals
        FITNESS-FN, a function that measures the fitness of an individual
```

repeat

new_population ← empty set

for i = 1 to SIZE(population) do

x ← RANDOM-SELECTION(population, FITNESS-FN)

y ← RANDOM-SELECTION(population, FITNESS-FN)

child ← REPRODUCE(x, y)

if (small random probability) then child ← MUTATE(child)

add child to new_population

population ← new_population

until some individual is fit enough, or enough time has elapsed

return the best individual in population, according to FITNESS-FN

```
function REPRODUCE(x, y) returns an individual
```

inputs: x, y, parent individuals

n ← LENGTH(x); c ← random number from 1 to n

return APPEND(SUBSTRING(x, 1, c), SUBSTRING(y, c + 1, n))

2. Game → zero-sum game, adversarial search problems

① Min-Max Tree

② Alpha-Beta Pruning: $O(CB^{\frac{m}{2}})$ (all adversarial game can use)
α: MAX's best option on path to root
β: MIN's best option on path to root

```
def max-value(state, α, β):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor, α, β))
        if v ≥ β return v
    α = max(α, v)
    return v
```

```
def min-value(state, α, β):
    initialize v = +∞
    for each successor of state:
        v = min(v, value(successor, α, β))
        if v ≤ α return v
    β = min(β, v)
    return v
```

④ Max Pruning leaf = $(b-1)(M-1)$ remain the left branch of each branch from root
reorder leaves, # possible values = # leaf node - max pruning node.

⑤ Exploit Max: cut from ∇ , cut K largest
cut from Δ : cut K smallest

④ Monte Carlo Tree Search.

while

$$S1 \text{ Tree traversal } UCB_1(c_n) = \frac{u(c_n)}{n(c_n)} + C \times \sqrt{\frac{\log N(PARENT(c_n))}{n(c_n)}}$$

S2 Node Expansion.

S3 Simulation (Rollout)

S4 Backpropagation

3 Logic

valid: true for all models

satisfiable: ∃ model. true

unsatisfiable: not true for all models

conjunctive normal form: $(P_1 \vee \dots \vee P_i) \wedge \dots \wedge (P_j \vee \dots \vee P_k)$

entail: $A \models B$.

① $A \models B$ iff $A \Rightarrow B$ is valid ② $A \models B$ iff $A \wedge \neg B$ is unsatisfiable.

Model checking

DPLL → is a variant of DFS

function DPLL-SATISFIABLE?(s) returns true or false sound (correct) & complete

inputs: s, a sentence in propositional logic

clauses ← the set of clauses in the CNF representation of s

symbols ← a list of the proposition symbols in s

return DPLL(clauses, symbols, {})

function DPLL(clauses, symbols, model) returns true or false

if every clause in clauses is true in model then return true

if some clause in clauses is false in model then return false

P, value ← FIND-PURE-SYMBOL(symbols, clauses, model) Pure Symbol Heuristic (all A or all B and assign as 1)

if P is non-null then return DPLL(clauses, symbols - P, model ∪ {P=value})

P, value ← FIND-UNIT-CLAUSE(clauses, model) Unit Clause Heuristic Eg. B ∨ F ∨ T ∨ F

if P is non-null then return DPLL(clauses, symbols - P, model ∪ {P=value}) only valid in clause

P ← FIRST(symbols); rest ← REST(symbols)

return DPLL(clauses, rest, model ∪ {P=true}) or

DPLL(clauses, rest, model ∪ {P=false}))

three rules of inference

① Modus Ponens: if knowledge base contain $A \wedge A \Rightarrow B$. we can infer B

② And-Elimination: if knowledge base contain $A \wedge B$. we can infer A and B.

③ Resolution: if knowledge base contain A and B. we can infer $A \wedge B$

Forward chaining KB?

function PL-FC-ENTAILS?(KB, q) returns true or false

inputs: KB, the knowledge base, a set of propositional definite clauses

q, the query, a proposition symbol

count ← a table, where count[c] is the number of symbols in c's premise

inferred ← a table, where inferred[s] is initially false for all symbols

agenda ← a queue of symbols, initially symbols known to be true in KB

while agenda is not empty do

p ← POP(agenda)

if p = q then return true

if inferred[p] = false then

inferred[p] ← true

for each clause c in KB where p is in c.PREMISE do

decrement count[c]

if count[c] = 0 then add c.CONCLUSION to agenda

return false

4 Bayes Net (tree-like net usually not holds ???)

mutually independent $(A \cup B)$, $P(A \cdot B) = P(A) \cdot P(B)$

conditional independent $B \perp\!\!\!\perp A | C$, $P(A \cdot B | C) = P(A | C) \cdot P(B | C)$, $P(C | A \cdot B) = P(C | A)$

$P(Q_1 \dots Q_k | E_1 \dots E_l) = \prod_i P(Q_i | E_1 \dots E_l)$: Q_i: Query variable E_i: Evidence variable, C: Hidden variable.

④ Infer by Enumeration

1. Collect all the rows consistent with the observed evidence variables.

2. Sum out (marginalize) all the hidden variables.

3. Normalize the table so that it is a probability distribution (i.e. values sum to 1)

Bayes Network Representation. CDAG 1

② Variable Elimination.

```

function ELIMINATION-ASK( $X, e, bn$ ) returns a distribution over  $X$ 
inputs:  $X$ , the query variable
         $e$ , observed values for variables  $E$ 
         $bn$ , a Bayesian network specifying joint distribution  $P(X_1, \dots, X_n)$ 
unnormalised  $factors \leftarrow []$ 
for each var in ORDER( $bn.VARS$ ) do
    factors  $\leftarrow [MAKE-FACTOR(var, e)]factors$ 
    if var is a hidden variable then factors  $\leftarrow SUM-OUT(var, factors)$ 
return NORMALIZE(POINTWISE-PRODUCT(factors))

```

order of elimination: best \rightarrow NP-hard problem:

Maybe good? ① least out-neighbour ② min-weight ③ min-fid: min size of factor

Approximate.

③ Prior Sampling 

④ Rejection Sampling: consistent with given info? reject it!

⑤ Likelihood Weight (evidence $\rightarrow P$, Query \rightarrow Sampling)

```

function LIKELIHOOD-WEIGHTING( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$ 
inputs:  $X$ , the query variable
         $e$ , observed values for variables  $E$ 
         $bn$ , a Bayesian network specifying joint distribution  $P(X_1, \dots, X_n)$ 
         $N$ , the total number of samples to be generated
local variables:  $W$ , a vector of weighted counts for each value of  $X$ , initially zero
for  $j = 1$  to  $N$  do
     $x, w \leftarrow WEIGHTED-SAMPLE(bn, e)$ 
     $W[x] \leftarrow W[x] + w$  where  $x$  is the value of  $X$  in  $x$ 
return NORMALIZE(W)

```

```

function WEIGHTED-SAMPLE( $bn, e$ ) returns an event and a weight

```

```

 $w \leftarrow 1; x \leftarrow$  an event with  $n$  elements initialized from  $e$ 
foreach variable  $X_i$  in  $X_1, \dots, X_n$  do
    if  $X_i$  is an evidence variable with value  $x_i$  in  $e$ 
        then  $w \leftarrow w \times P(X_i = x_i | parents(X_i))$ 
    else  $x[i] \leftarrow$  a random sample from  $P(X_i | parents(X_i))$ 
return  $x, w$ 

```

Gibbs Sampling (consistent)

```

function GIBBS-ASK( $X, e, bn, N$ ) returns an estimate of  $P(X|e)$ 
local variables:  $N$ , a vector of counts for each value of  $X$ , initially zero
 $Z$ , the non-evidence variables in  $bn$ 
 $x$ , the current state of the network, initially copied from  $e$ 
initialize  $x$  with random values for the variables in  $Z$ 
for  $j = 1$  to  $N$  do
    for each  $Z_i$  in  $Z$  do
        set the value of  $Z_i$  in  $x$  by sampling from  $P(Z_i | mb(Z_i))$ 
         $N[x] \leftarrow N[x] + 1$  where  $x$  is the value of  $X$  in  $x$ 
    return NORMALIZE(N)

```

*markov blanket - parents
child's parent*

Type of Environment in AI

- ① Fully observable / Partially Observable: sense the complete state at each point time
- ② static / dynamic: the environment keep change?
- ③ stochastic / deterministic: a uniqueness in the agent's current state determine next state
- ④ know physics: agent know the transitive model.